

# COMS W4701: Artificial Intelligence

## Lecture 2: Search Problems

Tony Dear, Ph.D.

Department of Computer Science  
School of Engineering and Applied Sciences

# Today

---

- Search problem formulation
- State space graphs and search trees
- Uninformed search: DFS, BFS, UCS
- Informed search: Greedy, A\*
- Search heuristics: Admissibility, design

# Problem-Solving Agents

---

- Goal-based agent that defines goals as a set of *world states*—descriptions of the task environment at the current time
- Assume our environments are **observable, discrete, deterministic, static**
  - Percepts are trivial, since we see entire environment
  - Action results always known, go from one state to another state
- We will find an *action sequence* that will result in a *state sequence* to a goal state
- This is the agent's solution to a **search problem**


# Search Problems

---

- **State space  $S$ :** Set of descriptions of the agent and environment
- **Actions:** (Finite) set of available actions in a state
  - Ex:  $Actions(s_1) = \{a_1, a_2, a_3\}$
- **Transition model:** Mapping from (state, action) to a new state
  - Ex:  $Result(s_1, a_1) = s_2$
- **Action costs:** Numerical cost for a (state, action, new state) transition
  - Ex:  $Cost(s_1, a_1, s_2) = 10$
- **Goal test (for goal states)**
  - Ex:  $IsGoal(s_1) = \text{False}$ ,  $IsGoal(s_2) = \text{True}$


# Example: Grid World Path Finding

- **State space:** Current coordinates of the rabbit
  - $S = \{(x, y) \mid 0 \leq x \leq 3, 0 \leq y \leq 4\}$
- **Actions:**  $Actions((x, y)) = \{\text{Up, Down, Left, Right}\}$
- **Costs:**  $Cost(s, a, s') = 1, \forall s, a, s'$
- **Transition model:**  $Result((x, y), \text{Up}) = (x, y + 1)$ ,  $Result((x, y), \text{Down}) = \dots$ 
  - Should also account for walls and boundaries, e.g.  $Result((0,0), \text{Left}) = (0,0)$
- **Goal test:**  $In((3,3))?$



# Multiple Carrots?

- What has changed about the problem?
- State space
  - Location of rabbit, Booleans indicating carrots eaten?
- Transition model
  - Update both rabbit location as well as carrot Boolean if locations match
- Goal test
  - Are all carrots eaten? Are all Boolean indicators True?



# Search Problem Example: $n$ -puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- **State:**
- **Action:**
- **Action cost:**
- **Goal test:**

# More Search Problem Examples


---

- Route-finding (e.g., vehicle navigation), robot navigation in the real world
- Touring problems (traveling salesperson)
- Layout and assembly sequencing problems
- Mathematical puzzles and proofs: Infinitely large state spaces!
- Knuth's conjecture (1964): Starting with the number 4, use a combination of factorial, floor, and sqrt operations to reach any other desired integer
- States: All nonnegative integers




# State Space Graphs

- **State space graph:** A mathematical representation of a search problem
  - *Vertices* are states; *edges* are actions
  - Each state occurs only once!
- *Paths* are sequences of actions/states
- A *solution* is a path from initial to goal state
- We can rarely build this full graph in memory—it can be very large or infinite



# Search Trees


- Need a systematic way of performing search over a state space graph
- **Search tree:** Nodes are states, edges are actions; root is initial state




- Unlike state space graph, states can occur more than once
- Each node corresponds to a *unique* path from initial state

# State Space Graphs vs. Search Trees


State Space Graph



Search Tree



# General Search Ideas



- From current node, **expand** and consider all possible actions
- Generate successor **nodes** for each resultant state according to transition function
  - Each node should track its corresponding state, parent, prior action, and total cost so far
- Successors are added to a **frontier** of possible next nodes to expand
- Frontier forms a boundary between explored and unexplored parts of tree

# Implementation Details

---

- What does *expansion* of children nodes entail?
- Find new state using transition function; store parent, action, and new *cumulative* cost
- How to *select* a node to expand from the frontier?
- **Best-first search:** Use an **evaluation function**  $f(n)$  assigning each node a priority
- **Uninformed search:**  $f(n)$  has no knowledge about how close a state is to goal
- What to do with states appearing more than once in search tree?
- Idea: Keep track of all *reached* states and costs in a lookup table
- A reached state should only be reconsidered if we find a cheaper path to it!



# Best-First Search

```
function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure
  node  $\leftarrow$  NODE(STATE=problem.INITIAL)
  frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
  reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    if problem.IS-GOAL(node.STATE) then return node
    for each child in EXPAND(problem, node) do
      s  $\leftarrow$  child.STATE
      if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
        reached[s]  $\leftarrow$  child
        add child to frontier
  return failure
```

```
function EXPAND(problem, node) yields nodes
  s  $\leftarrow$  node.STATE
  for each action in problem.ACTIONS(s) do
    s'  $\leftarrow$  problem.RESULT(s, action)
    cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
    yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```


# Depth-First Search

- Idea: Expand a *deepest* node first, implement frontier as a **stack** (LIFO)
- Behavior: Frontier expands toward tree leaves
- **Early goal test** can be done **when adding to** rather than *popping* from frontier




# DFS Properties

- **Time complexity:** How many nodes to explore in the worst case?  $O(b^m)$
- **Space complexity:** How many frontier nodes to keep in memory?  $O(bm)$
- **Completeness:** Guaranteed to find solution? Not if state space is infinite
- **Optimality:** Solution guaranteed to be lowest cost? No, only returns first solution




- $b$  is the *branching factor*
- $m$  is the *maximum depth*
- Total nodes:  $O(1 + b + b^2 + \dots + b^m)$







# BFS Properties

- **Time complexity:** How many nodes to explore in the worst case?  
 $O(b^d)$
- **Space complexity:** How many frontier nodes to keep in memory?  
 $O(b^d)$
- **Completeness:** Guaranteed to find solution? If solution exists, yes!
- **Optimality?** Solution guaranteed to be lowest cost? Only if costs are uniform




- $d$  is depth of the shallowest solution
- May be significantly smaller than  $m$
- Max frontier size is  $O(b^d)$

# BFS vs DFS





# Improving DFS and BFS

- **Depth-limited DFS:** Prevent DFS from going past a set depth  $l$
- Time complexity  $O(b^l)$ , space complexity  $O(bl)$
- Best if we know *diameter* of state space and check for short cycles
- **Iterative-deepening:** Iteratively do depth-limited search with increasing  $l$ : try  $l = 0$ , then  $l = 1, \dots$
- Ends when  $l$  reaches  $d$  (depth of shallowest solution)
- Time complexity  $O(b^d)$ , space complexity  $O(bd)$
- Why is wasted effort in upper levels of search tree not a concern?





# Uniform-Cost Search (Dijkstra)

- Idea: Expand node that least increases total cost, implement frontier as priority queue
- Evaluation function  $f(n)$  is the *total path cost so far*
- States may be added and replaced multiple times!
- Goal test must be done **when popping from frontier**, not when adding in





# BFS vs UCS



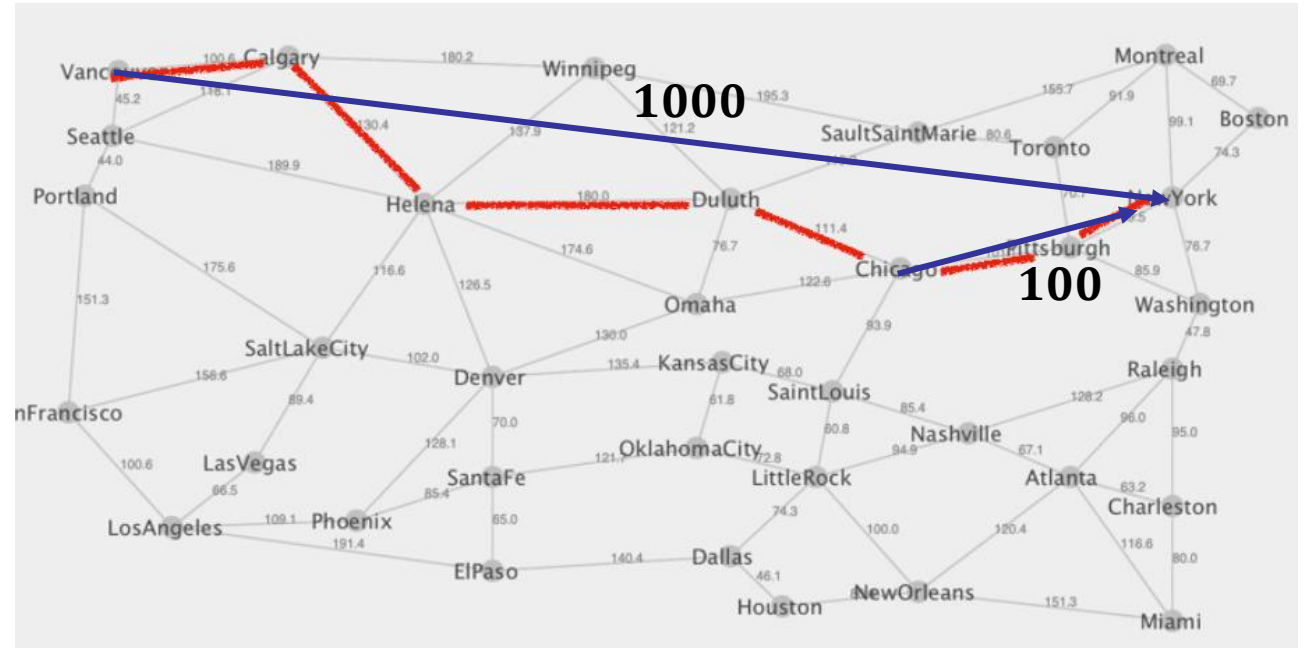
# UCS Properties

- Let  $C^*$  be the cost of optimal solution
- Let  $\epsilon$  be lower bound on all possible costs
- $1 + C^*/\epsilon$  is the max depth to traverse before finding optimal solution
- **Time and space complexity:  $O(b^{1+C^*/\epsilon})$**
- UCS is both **complete** and **optimal**



# Informed (Heuristic) Search

- Oftentimes we have additional, *domain-specific heuristics* that tell us how close a state is to a goal
- **Heuristic function  $h(n)$ :** Estimated cost of cheapest path from state at node  $n$  to a goal state
- Often come from *relaxed problems*, precomputed *subproblem* solutions, or learning from experience




Example: Euclidean or Manhattan distance on a map





# Greedy Best-First Search

- Idea: Expand node that *appears* closest to goal according to heuristic function
- Evaluation function is the heuristic function!  $f(n) = h(n)$
- Again implement frontier using priority queue




# Greedy Search Properties

- Performance depends entirely on usefulness of heuristic function
- Best case: Go straight toward the goal
- Worst case: Like a badly misguided DFS
- Complete in finite state spaces
- No guarantee of optimality since true costs are never considered







# A\* Search

- Idea: From a given node, estimate the *best* path that *continues* to the goal
- $f(n) = g(n) + h(n)$ : Sum of path cost to  $n$  and estimated cost from  $n$  to goal
- Benefits of both UCS and greedy best-first search





# A\* vs UCS vs BFS


- BFS expands search tree by increasing depth
- UCS expands acc. to increasing  $g$ -cost
- Contours are “circular” around start state, if normalized by path costs
- A\* expands acc. to increasing  $g + h$  cost
- If heuristic is good, expanded states should show preference toward goal



# A\* Examples



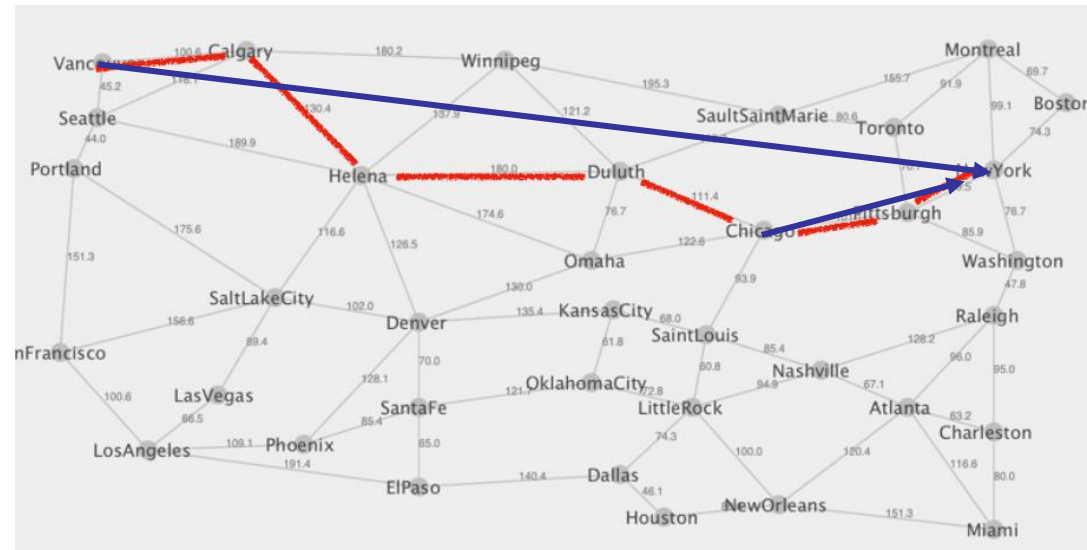
# When is A\* Optimal?



- What is the problem here?
- Heuristic along optimal path overestimated the true cost!
- **Good heuristics** should be optimistic—never overestimate true costs

# Admissible Heuristics


- A heuristic  $h$  is **admissible** if  $0 \leq h(n) \leq h^*(n)$  where  $h^*(n)$  is true cost from  $n$  to goal
- Most heuristics derived from relaxed problems are admissible
- Same state space graph, but with added edges
- With fewer constraints or restrictions, problems are easier to solve
- Example: Euclidean distances



# Example: 8-Puzzle

7	2	4
5		6
8	3	1

Start State



	1	2
3	4	5
6	7	8

Goal State



# Misplaced Tiles Heuristic

- $h(n)$  = number of misplaced tiles, not including blank

$$h\left(\begin{array}{|c|c|c|}\hline \blacksquare & 1 & 2 \\ \hline 3 & 4 & 5 \\ \hline 6 & 7 & 8 \\ \hline \end{array}\right) = 0 \quad h\left(\begin{array}{|c|c|c|}\hline 1 & 4 & 2 \\ \hline \blacksquare & 5 & 8 \\ \hline 3 & 6 & 7 \\ \hline \end{array}\right) = 7$$

- **Relaxed problem**: Any tile can be correctly replaced with just one move
- **Admissible** because misplaced tiles will always require *at least* one move

# Manhattan Distance Heuristic

- $h(n)$  = sum of Manhattan distances between current tile positions and goal positions

$$h(start) = 18$$

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State


- Relaxed problem: Multiple tiles can simultaneously occupy same space
- Admissible because misplaced tiles will always require *at least* number of moves equal to Manhattan distance

# Heuristic Domination

- For any node  $n$ , Manhattan distance heuristic  $h_2(n) >$  misplaced tiles heuristic  $h_1(n)$
- $h_2$  **dominates**  $h_1$  if  $h_2(n) \geq h_1(n)$  for all  $n$
- A\* search using  $h_2$  will be more efficient and never expand more nodes than  $h_1$
- $h_2$  reflects true costs more accurately
- Suppose we have collection of admissible heuristics  $h_1, h_2, \dots, h_m$
- The composite heuristic  $h(n) = \max\{h_1(n), \dots, h_m(n)\}$  is admissible and dominates all other heuristics!

# Completeness and Optimality of A\*

- A\* is complete (same reason as greedy, UCS, BFS)
  - If heuristic function is admissible, A\* is also optimal!
  - Proof by contradiction:
    - Assume A\* returns a suboptimal solution with cost  $C > C^*$
    - Then there exists some unexpanded node  $n$  on optimal path
    - Since  $n$  was not expanded,  $f(n) > C^*$
    - However,  $f(n) = g(n) + h(n) = g^*(n) + h(n) \leq g^*(n) + h^*(n) = C^*$  Contradiction!
- By definition      Since  $n$  is on optimal path      Since  $h$  is admissible      By definition



# Satisficing Solutions

- Like BFS or UCS, A\* may suffer computationally intractable memory requirements
- Idea: Trade off admissibility for more accurate heuristics to reduce computation
- Return **satisficing solutions**—suboptimal, but “good enough”
- **Weighted A\* search:**  $f(n) = g(n) + \alpha h(n)$
- We can choose to place higher weight  $\alpha$  on the heuristic
- Generalizes A\* ( $\alpha = 1$ ), UCS ( $\alpha = 0$ ), and greedy best-first ( $\alpha = \infty$ )
- Suboptimality: If optimal solution has cost  $C^*$ , weighted A\* solution may cost up to  $\alpha C^*$

# Memory-Bounded Search

---

- We can also consider  $A^*$  variants that are more memory-efficient
- **Beam search:** Limit frontier size by discarding worst nodes past a given limit
- Alternatively, discard nodes with scores much smaller than best one
- **Iterative-deepening  $A^*$  (IDA\*):** Repeatedly run  $A^*$  with increasing depth limit
- Nodes with higher  $f$ -cost than limit are treated as leaves
- Increment depth limit by smallest  $f$ -cost of “leaves” from previous iteration
- IDA\* worst case: Each node has different  $f$ -cost, num iterations equal to num states

# Summary

---

- Objective of search problems is to find action/state sequence to reach a goal state
- Represented by state space graphs; search algorithms follow a tree structure
- Uninformed search: No usage of information indicating closeness to goal
- Examples: Depth-first, breadth-first, depth-limited, iterative deepening, uniform-cost
- Generally suffer from lack of completeness or intractable memory usage
- Informed search: Domain-specific heuristics guide search toward goal
- Greedy best-first and A\* search use a heuristic function to evaluate frontier nodes
- Optimal if heuristics are admissible: good, optimistic estimates of true costs