

# COMS W4701: Artificial Intelligence

## Lecture 3: Constraint Satisfaction Problems

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences

# Today

---

- Constraint satisfaction problems
- Backtracking search
- Ordering heuristics
- Inference and constraint propagation
- Local consistency

# States with Structure

---

- Task environments so far: Fully observable, discrete, deterministic, static
- Transitions and heuristics are problem-specific
- States are *atomic* black boxes
- **Planning** solution is an action **sequence**, or state space **path**

Using a set of variables (features) to represent a domain.

- What if our states have a common **factored representation**?
- We can use *general-purpose* heuristics
- Solution of an **assignment** problem is the **goal** itself, not a path


# Constraint Satisfaction Problems

---

- Special structured search problems with 3 components
  - Variables  $X = \{X_1, \dots, X_n\}$
  - Domains  $D = \{D_1, \dots, D_n\}$
  - Constraints  $C = \{C_1, \dots, C_m\}$
- Goal test: A **complete, consistent** assignment of values to each variable  $X_i$  from respective domain  $D_i$  s.t. all constraints  $C$  are satisfied
- We may have incomplete or inconsistent assignments along the way


# Example: Map Coloring

- Goal: Color a map so that no adjacent territories have the same color
- Variables:  $X = \{WA, NT, Q, NSW, V, SA, T\}$
- Domains:  $D_i = \{\text{red, green, blue}\}$
- Constraints: Implicit vs explicit representation
  - $C = \{WA \neq NT, WA \neq SA, NT \neq SA, NT \neq Q, SA \neq Q, Q \neq NSW, NSW \neq V\}$
  - $C = \{(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), \dots\}, \dots\}$



# Constraint Graphs

- Visualization of a CSP in a graph
- Nodes: Variables and domains
- Edges: Presence of binary constraints
- Certain graph algorithms can help give insight into the CSP
- E.g., a  $k$ -connected component indicates a  $k$ -nary constraint
- CSPs that are actually tree structures can be solved without backtracking




# Types of Constraints

---

- So far we've seen **binary** constraints relating two variables
- **Unary** constraints involve a single variable, equivalent to *domain reduction*
  - Ex:  $SA \neq \text{green}$  (implicit),  $SA \in \{\text{green}, \text{red}\}$  (explicit)
- Higher-order (**global**) constraints relate arbitrary number of variables
  - Ex: *Alldiff* requiring all variables to have different values
- If domains are finite, always possible to rewrite global constraints as binary ones using *auxiliary variables*

# Example: $n$ -Queens

- Place  $n$  queens on  $n \times n$  board s.t. none share a row, column, or diagonal
- Variables:  $X_{ij}$ ,  $1 \leq i \leq n$ ,  $1 \leq j \leq n$  (grid spaces)
- Domains:  $D_{ij} = \{0,1\}$  (queen or no queen)




- Constraints: 
$$\forall i \sum_j X_{ij} = 1 \quad \forall i, j \sum_k X_{i+k, j+k} \leq 1$$
$$\forall j \sum_i X_{ij} = 1 \quad \forall i, j \sum_k X_{i+k, j-k} \leq 1$$



# Example: $n$ -Queens

- Alternatively, states can just represent a row of the board (vs grid space)
- Variables:  $X_i, 1 \leq i \leq n$  (row of the board)
- Domains:  $D_i = \{1, \dots, n\}$  (column containing queen)
- Constraints:  $\forall i, j, X_i \neq X_j$   
 $\forall i, j, X_j - X_i \neq |j - i|$



# Example: Cryptarithmic

- Variables:  $X = \{T, W, O, F, U, R, \dots\}$
- Also carry-overs!  $C_1, C_2, C_3$

$$\begin{array}{r} T \ W \ O \\ + \ T \ W \ O \\ \hline F \ O \ U \ R \end{array}$$

- Domains:  $D_i = \{0, \dots, 9\}$

- Constraints

- $Alldiff(X_i)$ 

$$\begin{aligned} O + O &= 10C_1 + R \\ W + W + C_1 &= 10C_2 + U \\ T + T + C_2 &= 10F + O \\ C_3 &= F \end{aligned}$$

- 7 possible solutions!

$$734 + 734 = 1468$$

$$765 + 765 = 1530$$

$$836 + 836 = 1672$$

$$846 + 846 = 1692$$

$$867 + 867 = 1734$$

$$928 + 928 = 1856$$

$$938 + 938 = 1876$$

# Example: Sudoku

- Variables: One for each open square
- Domains:  $\{1, \dots, 9\}$
- Constraints:
  - *Alldiff* for each column
  - *Alldiff* for each row
  - *Alldiff* for each  $3 \times 3$  square

					8			4
	8	4		1	6			
			5			1		
1		3	8			9		
6		8				4		3
		2			9	5		1
		7			2			
			7	8		2	6	
2			3					

# Solving CSPs with Search

---

- **States:** Partial assignments (initial state is no assignment)
- **Actions:** Assign value to an unassigned variable from its domain
- **Goal test:** Complete, consistent assignment
- No explicit costs, so maybe something like DFS/BFS?
- Problem: Naïve implementation -> lots of repeated states!
- **Branching factor:** Number of unassigned variables  $\times$  size of domain
  - Branching factor at root:  $n \times d$  for  $n$  variables with domain size  $d$
  - Branching factor at 2<sup>nd</sup> level is  $(n - 1)d$ , then  $(n - 2)d$  at 3<sup>rd</sup> level, ...
  - Total number of leaves / possible goal states:  $O(n! \times d^n)$

# Backtracking Search


---

- Idea: CSPs are *commutative*, order of variable assignment doesn't matter
  - (WA = red, NT = green) is the same as (NT = green, WA = red)
- Each search tree tier only needs to correspond to a single variable
  - Branching factor is  $d$  at every level, number of leaves is  $O(d^n)$
- Running DFS allows us to make one variable assignment at a time and *backtrack* by undoing inconsistent assignments to try alternatives
  - Don't need a frontier—just keep track of what domain values are available



# Variable Selection


- Heuristic for picking next variable: Minimum remaining values (MRV)
  - Also called “most constrained variable” (MCV) or “fail-first”
  - Choose the variable with the fewest legal left values in its domain



- What if we have multiple MRV variables (e.g., initial assignment)?
  - **Degree** heuristic: Pick variable that appears in the most constraints
  - Reduce branching factor on future choices

# Domain Value Assignment

- Heuristic for assigning a value: Least constraining value (LCV)
  - Choose a value that imposes fewest constraints on future assignments
  - May require some computation
- Why least, not most?
  - We don't have to use all values, just find an assignment that works
  - Try to look for the most likely ones earlier
  - Keep more options open for other variables





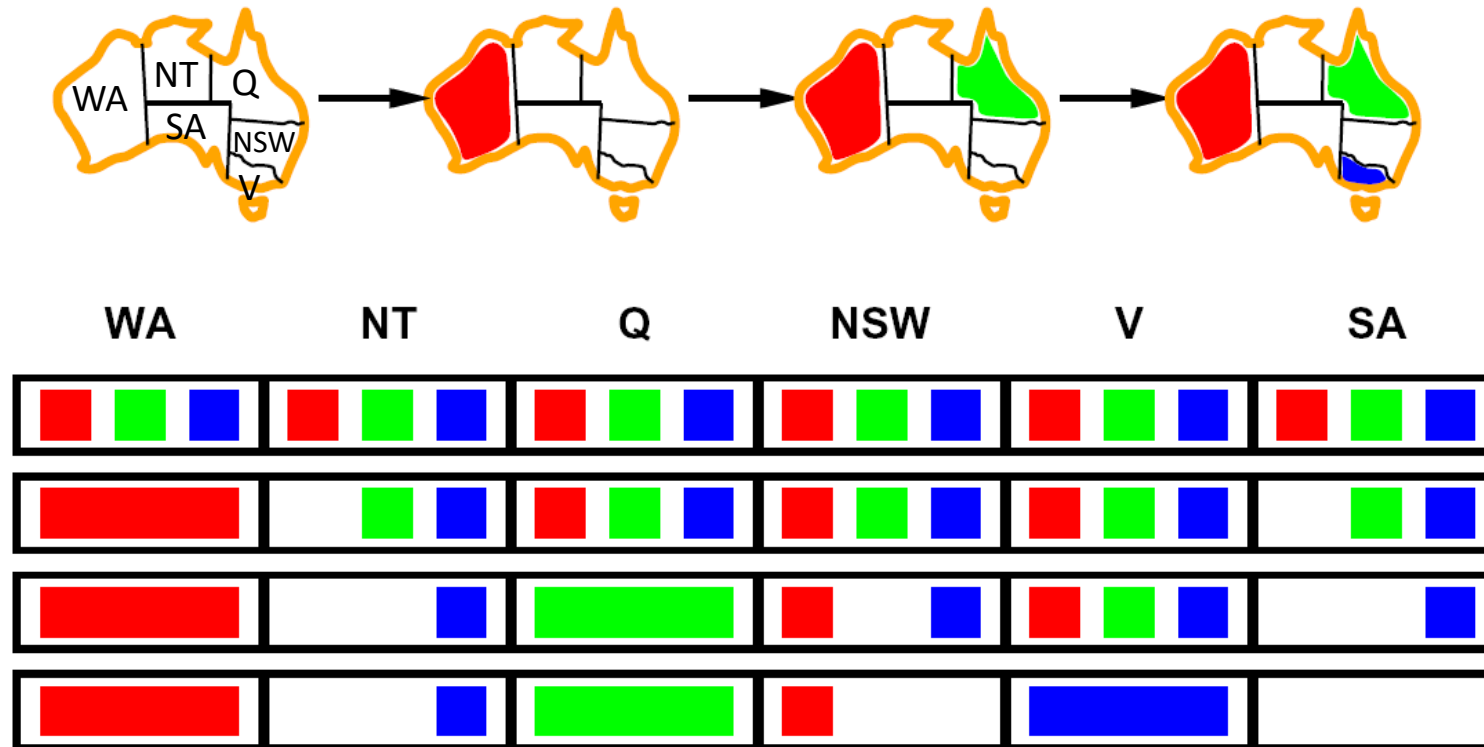
# Improving Backtracking: Inference

---

- Both MRV and LCV heuristics assume some sort of “forward checking”
- We can detect inevitable failure earlier by checking constraints as we go
- Idea: Delete inconsistent domain values as we make assignments
- If any domain is left empty, we should not proceed
- Domain reduction also better informs MRV on selecting variables
- Inference can also inform how to use the LCV heuristic

# Forward Checking

- **Forward checking:** After assigning a variable, check other variables related to it by a constraint and eliminate inconsistent domain values
- If any domain becomes empty, current assignment is denoted a failure



The image displays a 4x6 grid of colored squares, where each square represents a specific operation in a sequence. The colors used are red, green, and blue. The grid is organized into four rows and six columns. The first row contains six small squares, each with a single color: red, green, blue, red, green, blue. The second row contains six larger squares, each with a single color: red, green, blue, red, green, blue. The third row contains six larger squares, each with a single color: red, green, blue, red, green, blue. The fourth row contains six larger squares, each with a single color: red, green, blue, red, green, blue.

# Local Consistency


- Forward checking is an example of **constraint propagation**
- Can be interleaved with search, or done as a preprocessing step prior to search
- Result is **local consistency**: remaining domain values do not violate any constraints
- **Node consistency**: All unary constraints are satisfied
  - Can always be done prior to starting any search process
- **Arc consistency**: No binary constraint is violated by any variable's domain values
  - $X_i$  is arc-consistent with respect to  $X_j$  if every value in  $D_i$  can be paired with a value in  $D_j$  such that no constraint between  $X_i$  and  $X_j$  is violated.

# Arc Consistency

- $X_i$  being arc consistent with  $X_j$  does not imply that  $X_j$  is arc consistent with  $X_i$ !
- Example:  $Y = X^2$ , both with domains  $\{0,1, \dots, 9\}$ 
  - To make  $X$  arc consistent with  $Y$ : Reduce domain of  $X$  to  $\{0,1,2,3\}$
  - To make  $Y$  arc consistent with  $X$ : Reduce domain of  $Y$  to  $\{0,1,4,9\}$
- Entire CSP is arc consistent iff both domains are reduced
- Another issue: Arc consistency may not be preserved if a domain is reduced
- If  $X_i$  is arc consistent with  $X_j$  and  $X_j$  is changed (e.g. due to forward checking),  $X_i$  *may no longer be arc consistent with  $X_j$*



# AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check





# AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check





# AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check



# AC-3 Algorithm



- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check






# AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check




# AC-3 Algorithm

- To enforce arc consistency, e.g. during backtracking search, *loop over all constraints* whose variables' domains have changed until all are arc consistent
- More general than forward checking
- Must maintain a dynamic queue of constraints to check



# Path and $k$ -Consistency

- At least one arc inconsistent  $\Rightarrow$  no solution
- BUT arc consistency not sufficient to guarantee a solution!
- Idea: Look at more than two variables at a time



- $\{X_i, X_j\}$  is **path consistent** with  $X_m$  if for every arc consistent assignment to  $\{X_i, X_j\}$ , there is an assignment to  $X_m$  consistent with  $X_i$  and  $X_j$
- **$k$ -consistency**: For every consistent assignment to a set of  $k - 1$  variables, a consistent value can be assigned to the  $k$ th variable
- A **strongly  $n$ -consistent** CSP ( $n, n - 1, \dots, 1$  consistent) is guaranteed a solution, but determining  $n$ -consistency is exponentially hard (NP-complete)!

# Special Constraints


- Consistency for certain constraint types can be checked quickly
- *Alldiff* with  $m$  variables and  $n$  unique values: immediately declare no solution if  $m > n$
- Suppose we have domains with lower and upper bounds
- **Bounds propagation**: If constraints are equality or inequality (e.g., **resource constraints**), we can use them to tighten the bounds and make them consistent
- Ex:  $X_1, X_2, X_3$ , all with domains  $[1,5]$ 
  - Constraint  $\sum X_i = 13$ : Reduce domains to  $[3,5]$
  - Constraint  $\sum X_i \leq 5$ : Reduce domains to  $[1,3]$

# Introducing New Constraints

- We can also improve CSP solvers by *introducing new constraints*
- Recall: We backtrack whenever we see an inconsistent solution, either in the present (leaf of the search tree) or in the future (via constraint propagation)
- **Constraint learning**: Record the current assignment as a constraint
- A CSP may have multiple solutions due to *value symmetry*
- Ex: In map coloring, there are  $d!$  solutions by permutation
- Idea: Introduce **symmetry-breaking constraints** (e.g.,  $NT < SA < WA$ ) to reduce the number of solutions and shrink the search tree

# Local Search

- Methods so far: Build up solution incrementally, check constraints
- **Local search**: Start with an arbitrary complete assignment, *modify* it until consistent
- No frontier to maintain, no backtracking!




- **Min-conflicts**: Reassign a variable to a value that *minimizes conflicts*

**function** MIN-CONFLICTS(*csp*, *max\_steps*) **returns** a solution or *failure*  
**inputs:** *csp*, a constraint satisfaction problem  
          *max\_steps*, the number of steps allowed before giving up

*current*  $\leftarrow$  an initial complete assignment for *csp*  
**for** *i* = 1 to *max\_steps* **do**  
    **if** *current* is a solution for *csp* **then return** *current*  
    *var*  $\leftarrow$  a randomly chosen conflicted variable from *csp*.VARIABLES  
    *value*  $\leftarrow$  the value *v* for *var* that minimizes CONFLICTS(*csp*, *var*, *v*, *current*)  
    set *var* = *value* in *current*  
**return** *failure*


# Example: 4-Queens



- Can be very effective in practice, roughly independent of problem size
- E.g., can solve the 1M-queens problem in an average of 50 iterations
- Also useful for solving problems online as they change
- For example, repair existing solutions as new constraints come and go

# Local Search and Optimization\*

- **Local search** generally useful when we care about goal more than path
- E.g., optimization problems with **objective function**, no specific goal test
- Pros: Very little memory, can work well in large or infinite state spaces






# Hill-Climbing Methods\*

---


- Many local search algorithms fall into category of *hill-climbing*
- No search tree, no knowledge beyond immediate neighbors
- Greedy search method, can get stuck at local maxima if unlucky
  
- Variations of simply ascending the hill:
  - **Stochastic**: Pick random moves from time to time
  - **First-choice**: Pick a random move that leads to a better successor
  - **Random-restart**: Try different initial states if stuck

# Other Methods\*

- **Simulated annealing:** Combine hill climbing with random walks to get unstuck from local maxima
- **Local beam search:** Maintain parallel searches but share information among threads
- **Genetic algorithms:** Generate successor states by combining two parent states



<https://eat.zesty.com/post/how-simulated-annealing-can-improve-your-lunch>



# Summary

---

- CSPs are assignment search problems with very specific structure
- Unlike general DFS, backtracking search can be very effective
- General-purpose heuristics: MRV, LCV
- Inference can further improve search performance
- Maintain local consistency through constraint propagation, take advantage of or add new common constraints
- Local search: Alternative search method with few memory requirements