

COMS W4701: Artificial Intelligence

Lecture 5: Sequential Decision Problems

Tony Dear, Ph.D.

Department of Computer Science

School of Engineering and Applied Sciences


Today

- Multi-armed bandit problems
- Exploration vs exploitation
- UCB action selection

- Markov decision processes
- Utilities, discounting, values, and policies
- Bellman equations

Multi-Armed Bandits

- Thought experiment: Suppose we have several slot machines
- Each machine has different rewards and different odds
- We can only find out by trying different machines
- Tradeoff between **exploration** and **exploitation**
 - Gather more information or maximize best rewards so far?
 - How to determine when current knowledge is good enough?
- Applications: Resource allocation for maximizing productivity, clinical trials to explore different treatments, financial portfolio design



Action Values

- Choosing a slot machine is choosing an *action* a
- A_t is the action chosen and R_t is the resultant reward at time t
- Action value is the *expected* reward of an action: $Q^*(a) = E[R_t | A_t = a]$
- If we know all Q^* , optimal strategy would be to pick a with highest Q^*
- Idea: Build *estimates* of Q^* by trying different actions and recording results
- First initialize all $Q_0(a)$, e.g. by trying each action once and recording reward
- *Sample-averaging*: $Q_t(a) = \frac{\text{sum of rewards from taking } a \text{ prior to } t}{\text{number of times taking } a \text{ prior to } t}$

Action Selection

- How to select actions to try while building up Q estimates?
- *Greedy* action selection always exploits: $A_t = \operatorname{argmax}_a Q_t(a)$
- Problem: We would never explore new actions
- What if there are better options out there?
- **ϵ -greedy**: Select greedy action *most* of the time, but with small probability ϵ , pick a random action to *explore* instead
- In the limit, estimates Q_t will converge to true Q^*

Controlling Exploration

- While exploration is *necessary* to learn action values, it is usually **not** optimal
- Too much exploration will lower total rewards received
- Example: Bandits with deterministic rewards only require one round of exploration
- Even for stochastic problems, we probably don't need to explore *forever*
- Better idea: Explore more beginning (higher ϵ), then gradually taper off (lower ϵ)
- Another problem: When exploring, ϵ -greedy picks actions completely randomly
- Can we do targeted exploration, e.g. pick actions that are more promising?

Upper Confidence Bound

- Idea: Target actions that have high *potential* for being optimal
- The *confidence interval* of a value estimate specifies range of true values
- Interval (estimate uncertainty) shrinks as we try an action more often
- **Upper confidence bound** action selection: $A_t = \operatorname{argmax}_a [Q_t(a) + \underbrace{c\sqrt{\ln t / N_t(a)}}_{UCB_t(a)}]$
 - $N_t(a)$ = number of times a has been taken prior to t
 - $\sqrt{1/N_t(a)}$ proportional to standard deviation
 - c is a parameter that controls the amount of exploration
- As $N_t(a)$ increases, $UCB_t(a)$ approaches $Q_t(a)$ due to increased confidence

Sequential Decision Problems



- Suppose we require a *sequence* of decisions in a *stochastic* environment
- Like stochastic games, we will have a *stochastic* transition model
- Unlike in games, rewards can be received at any time (not just the end)
- The problems will be single-agent, sequential, stochastic, fully observable
- A step-by-step *plan* no longer works due to stochasticity
- Instead, we form a *policy*: a mapping from states to actions

Markov Decision Processes


- Bandit problems are decision problems without states
- In **Markov decision processes**, agents take actions, move from state to state according to a transition function, and receive rewards along the way
- Set of states S and set of actions A
 - S may include initial and/or terminal states stochastic, so we need 3D matrix
- **Transition function** $T: S \times A \times S \rightarrow [0,1]$, where $T(s, a, s') = \Pr(s'|s, a)$
 - Also called the *model* or dynamics of the problem
- **Reward function** $R: S \times A \times S \rightarrow \mathbb{R}$, written as $R(s, a, s')$
- MDPs are **Markovian** in that we only need to track current, not past, states
The future is independent of the past given the present

Example: Gridworld

- States: Grid locations (11 in total)
 - First two states in column 4 are **terminal states**
- Actions: North, south, east, west (4 for each state)
 - May stay in original state if moving into wall
- Transition function: *Intended* direction with prob 0.8; otherwise, slip left or right with prob 0.1, respectively
- Rewards: +1 or -1 for moving into terminal states; small negative reward otherwise (*living reward*)



Example: Race Car MDP



MDPs in Practice

- Agriculture
 - S: Soil condition and precipitation forecast. A: Whether or not to plant a given area.
- Water resources and energy generation
 - S: Water levels and inflow. A: How much water to use to generate power.
- Inspection and maintenance
 - S: System age and probability failure. A: Whether to test / restore / repair a system.
- Inventory
 - S: Inventory levels and commodity prices. A: How much to purchase.
- Finance and investment
 - S: Holding or capital levels. A: How much to invest.
- Many, many more (D. J. White 1993)

Utilities

- A sequence of states and actions can be quantified by its **utility**
- *Rational* agent seeks to maximize its utility over time
- **Finite-horizon** MDPs: Process ends after some finite time T
 - Equivalent to entering a *terminal state* s_T
 - One way to define utility of a state-action sequence: sum up individual rewards

$$V_h([s_0, a_0, s_1, a_1, \dots, a_{T-1}, s_T]) = \sum_{t=0}^{T-1} R(s_t, a_t, s_{t+1})$$

- This won't work for **infinite-horizon** MDPs!

Discounted Utilities

- Which reward sequence is better? $R_1 = (1,1,1)$ vs $R_2 = (0,0,3)$
- Sums are equal, but R_1 is preferable if rewards *now* are better than rewards *later*
- Supported by psychology, economics; also due to uncertainty in obtaining rewards
- Idea: Apply a **discount factor** $0 < \gamma < 1$ to diminish future rewards
- Utility is computed using **additive discounted rewards**

$$V_h([s_0, a_0, s_1, a_1, \dots, a_{T-1}, s_T]) = \sum_{t=0}^{T-1} \gamma^t R(s_t, a_t, s_{t+1})$$

Discounting for Infinite-Horizon MDPs

- Additive utilities for finite-horizon MDPs use $\gamma = 1$
- For infinite-horizon MDPs, $\gamma < 1$ allows utility sequences to converge
 - The closer γ is to 0, the less we value the future
 - Rewards far in the future are effectively zero

$$V_h([s_0, a_0, s_1, \dots]) = \sum_{t=0}^{\infty} \gamma^t R(s_t, a_t, s_{t+1}) \leq \frac{R_{\max}}{1 - \gamma}$$

- Ex: Utility of infinite sequence of +1 rewards with $\gamma = 0.8$:

$$V_h = 1 + (0.8)1 + (0.8)^2 1 + \dots = \frac{1}{1 - 0.8} = 5$$

Policies and Value Functions



- Solving MDP means finding a **policy**—mapping from states to actions
- $\pi: S \rightarrow A$ tells agent what to do in any state
- Policies can be quantified by **value functions**
- $V^\pi: S \rightarrow \mathbb{R}$ is *expected utility of following π* starting from given state

$$V^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(s_t), s_{t+1}) \right], s_0 = s$$

- **Optimal** policy and value function:


$$\pi^* = \operatorname{argmax}_{\pi} V^\pi \quad V^* = \max_{\pi} V^\pi$$

Gridworld Policy and Value Function




$R(s) = -0.04$ for nonterminal states
 $\gamma = 1$ (no discounting)


Rewards and Policies




$$r < -1.6497$$



$$-0.7311 < r < -0.4526$$



$$-0.0274 < r < 0$$



$$r > 0$$


Recursive Relationship

$$V^\pi(s) = E \left[\sum_{t=0} \gamma^t R(s_t, \pi(s_t), s_{t+1}) \right]$$

- This is a function on every state in the state space
- For a given state s , we can alternatively write $V^\pi(s)$ as a *recursive function of successor state values*

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$


- $V^\pi(s)$ is a *weighted average* of (immediate reward plus discounted successor state values)



Solving for Values

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

- Suppose we know the model (T, R) , discount γ , and a fixed policy π
- Then the above is a system of $|S|$ *linear* equations in the $|S|$ unknowns $V^\pi(s)$
- Linear solvers: $\sim O(|S|^3)$ time, can find $V^\pi(s)$



Example: Mini-Gridworld

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

- States: A, B, C ; actions: left, right; $\gamma = 0.5$
- Policy: $\pi(s) = \text{left } \forall s$
- Rewards: $R(s, a, A) = +3, R(s, a, B) = -2, R(s, a, C) = +1$
- Transitions: $\text{Pr}(\text{intended direction}) = 0.8, \text{Pr}(\text{opposite direction}) = 0.2$
- V^π can be found by solving 3 linear equations:

+3	-2	+1
A	B	C

$$V^\pi(A) = 0.8(3 + 0.5V^\pi(A)) + 0.2(-2 + 0.5V^\pi(B))$$

$$V^\pi(B) = 0.8(3 + 0.5V^\pi(A)) + 0.2(1 + 0.5V^\pi(C))$$

$$V^\pi(C) = 0.8(-2 + 0.5V^\pi(B)) + 0.2(1 + 0.5V^\pi(C))$$

Bellman Optimality Equations

- Our goal is to find an **optimal policy** or **optimal value function**

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$




$$V^*(s) = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

$$\pi^* = \operatorname{argmax}_\pi V^\pi$$

$$V^* = \max_\pi V^\pi$$

- Bellman optimality equations are nonlinear!**




Value Function to Policy

- We don't know (yet) how to solve for V^* from scratch
- We *do* know how to find V^* given π^* (solve linear system)

- Bellman equation tells us how to find π^* given V^*

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Everything on the RHS is known!
- Solving for complete policy takes $O(|S|^2|A|)$ time



Example: Mini-Gridworld

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

+3	-2	+1
----	----	----

- States: A, B, C ; actions: left, right
- Transitions and rewards same as before

- Given $V^*(A) = 4.06, V^*(B) = 4.36, V^*(C) = 1.39$

- Find $\pi^*(B)$:

$$\gamma = 0.5 \quad \pi^*(B) = \operatorname{argmax} \begin{cases} 0.8(3 + 0.5V^*(A)) + 0.2(1 + 0.5V^*(C)) & \text{Left} \\ 0.8(1 + 0.5V^*(C)) + 0.2(3 + 0.5V^*(A)) & \text{Right} \end{cases}$$

Summary

- In sequential decision problems, an agent must take actions over and over
- Multi-armed bandits: No states, multiple actions; agent must often contend with exploration vs exploitation
- More general decision problems: Markov decision processes
 - States, actions, transitions, rewards, (discount factor)
- Goals -> rewards -> utilities -> value functions and policies
- Bellman optimality equations: Recursive and nonlinear